

Dynamic Programming in Futhark

Martin Elsman, DIKU, University of Copenhagen

October 23, 2021

This note describes how the Futhark [1, 2] dynamic programming library `dpsolver`¹ can be used to find fixpoints for functions from R^n to R^n and how solutions to multiple instances of a dynamic programming problem can be computed in parallel on GPUs. We give both single-dimensional and multi-dimensional examples and we show how the Futhark automatic differentiation feature may relieve programmers from specifying explicitly the Jacobian matrices, which are necessary for using `dpsolver`'s fast converging Newtonian functionality.

Introduction

A standard approach for finding fixpoints for numerical functions from R^n to R^n is to use the technique of successive approximations. Following Section 4 of Numerical Dynamic Programming in Economics, by John Rust [3], the dynamic programming solver that we shall apply here first uses a number of successive approximation steps before it applies a more efficient Newtonian method for narrowing in on a fixpoint. The latter method requires that the user specifies how to compute the Jacobian matrix (of type $R^{n \times n}$) given an approximate fixpoint. The Jacobian is then computed for each Newtonian step.

Example: Intersection of a circle and a quadratic equation

Following the example in Jim Lander's MAT 461/561 lecture notes, we first set out to find the intersection between the unit circle ($x_1^2 + x_2^2 = 1$) and the quadratic equation $x_2 = x_1^2$.²

We first define an operator for which we want to find a fixpoint. To ensure that the natural matrix norm of the Jacobian matrix for the function is less than 1 ($0 \leq x_1 \leq 1$ and $0 \leq x_2 \leq 1$), we give the following definition of the fixpoint operator G :

$$G(x_1, x_2) = (\sqrt{x_2}, \sqrt{1 - x_1^2})$$

Without having to define the Jacobian matrix for the function, we can find an approximation to the fixpoint using the successive approximation functionality of the `dpsolver` library.

¹The Futhark library `dpsolver` is based on a Matlab library implemented by Bertel Schjerning, ECON, University of Copenhagen.

²Analytically, the solution can easily be found by solving the quadratic equation $x_1^4 + x_1^2 - 1 = 0$, which leads to the solution $x_1 = 0.786151377757$ and $x_2 = 0.61803398874989$.

We first `import` the library `dpsolver` and instantiate the contained module `dpsolver` to the `f64` representation of floats:

```
import "dpsolver"
module dps = dpsolver f64
```

The function `dps.sa` that we shall apply has the following type:

```
val sa [m] : (f:[m]t->[m]t) -> (v:[m]t) -> (p:param) -> (b:t)
          -> ([m]t, bool, i64, t, t)
```

Here `t` is identical to `f64` due to the `f64` module instantiation of the `dpsolver` parameterised module.

We now define the `bellman` equation for which we want to find a fixpoint:

```
let bellman (x1:f64) (x2:f64) : (f64,f64) =
  (f64.sqrt x2, f64.sqrt(1 - x1**2))
```

Before we can make use of the library module, we define a small utility function `wrap2`, which takes a function of type `f64 -> f64 -> (f64, f64)` and turns it into a function of type `[2]f64 -> [2]f64`, which is compatible with the successive approximation functionality in the `dps` module:

```
let wrap2 (f: f64 -> f64 -> (f64,f64)) (a:[2]f64) : [2]f64 =
  let (x,y) = f a[0] a[1]
  in [x,y]
```

The following Futhark entry point makes a call to the `dps.sa` function with the above `bellman` function wrapped as a parameter:

```
entry test_sa (sa_max:i64) (sa_tol:f64)
  : ([2]f64, bool, i64, f64) =
  let v0 = [0.5, 0.5]
  let ap = dps.default with sa_max = sa_max
                      with sa_tol = sa_tol
  let (res, b, i, tol, _) = dps.sa (wrap2 bellman) v0 ap 0
  in (res, b, i, tol)
```

The function `dps.sa` also takes an initial approximation as argument (`v0`) together with an `ap` value that defines some slightly modified default parameter settings (max iterations, max tolerance, etc.)

We can now call the function:

```
> test_sa 60i64 1.0e-3f64
([0.7855137639650378f64, 0.6177294754734614f64], true, 56i64,
 8.769119278559945e-4f64)
```

We see that after 56 iterations, a fixpoint is found with a tolerance below `1e-3`, meaning that the last iteration step contributed to a change in value of less than `1e-3` for both x_1 and x_2 . For improved precision, many more iterations are required:

```
> test_sa 200i64 1.0e-9f64
([0.7861513784203592f64, 0.6180339890667096f64], true, 186i64,
 9.120002530949023e-10f64)
```

Faster convergence with Newton's method

The function G , as defined in (), has the following Jacobian matrix:

$$J_G(x_1, x_2) = \begin{bmatrix} 0 & \frac{1}{2\sqrt{x_2}} \\ \frac{-x_1}{\sqrt{1-x_1^2}} & 0 \end{bmatrix}$$

The following version of the `bellman` function takes its input as an array of size 2 and returns, along with the function result, the Jacobian matrix, relative to the argument:

```
let bellman_j (a:[2]f64) : ([2]f64, [2] [2]f64) =
  let x1 = a[0]
  let x2 = a[1]
  let res = [f64.sqrt x2, f64.sqrt(1-x1**2)]
  let j = [[0, 1/(2*f64.sqrt x2)],
            [-x1/(f64.sqrt(1-x1**2)), 0]]
  in (res, j)
```

The function `dps.poly` that we shall apply has the following type:

```
val poly [m]: (f: [m]t -> ([m]t,[m] [m]t)) -> (v:[m]t) -> (p:param)
  -> (b:t) -> ([m]t,[m] [m]t,bool,i64,i64,i64,t)
```

Again, here `t` is identical to `f64` due to the `f64` module instantiation of the `dpsolver` parameterised module. The function finds a fixpoint for the function `f` using a combination of successive approximation iterations and Newton-Kantorovich iterations. The initial guess is `v` and the parameter `p` controls the iteration passes. The function `f` should return a pair of a new next approximation and the Jacobian matrix for the function `f` relative to the argument given. The function returns a 7-tuple containing an approximate fixpoint, a Jacobian matrix for the fixpoint, a boolean specifying whether the algorithm converged (according to the values in `p`), the number of iterations used for the total `sa` iterations, the total Newton-Kantorovich iterations, and the number of roundtrips. The 7'th element of the result tuple is the tolerance of the last two fixpoint approximations (maximum of each dimension).

```
entry test_poly (sa_max:i64): ([2]f64, bool, i64, i64, i64, f64) =
  let v0 = [0.5, 0.5]
  let ap = dps.default with sa_max = sa_max
  let (res, _, b, i, j, k, tol) = dps.poly bellman_j v0 ap 0
  in (res, b, i, j, k, tol)
```

```
> test_poly 5i64
([0.7861513777574233f64, 0.6180339887498949f64], true, 5i64, 4i64, 1i64,
 1.1102230246251565e-16f64)
```

Notice that the programmer has manually provided code for computing the Jacobian matrix for the function. The result is a fixpoint with a tolerance below 1e-15, computed with an initial number of 5 successive approximation iterations followed by 4 Newtonial iterations (1 roundtrip was used).

Futhark AD

We can relieve the programmer from manually providing the code for the Jacobian matrix by using the automatic differentiation feature of Futhark, which provides a function `jvp` that performs forward automatic differentiation on arbitrary Futhark functions (featured in the Futhark clean-ad branch, but not yet featured in the `master` branch). An alternative is to encode float computations using so-called dual-numbers, following the approach of AD with dual numbers, but we shall not dive into this possibility here.

We first define a function `wrapj` that takes a function of type `[n] [m] . [n] f64 -> [m] f64` and turns it into a function of type `[n] [m] . [n] f64 -> ([m] f64, [m] [n] f64)` that, besides from the function result, returns the Jacobian matrix of the function:

```
let idd n i = tabulate n (\j -> if i==j then 1f64 else 0f64)

let wrapj [n] [m] (f: [n] f64 -> [m] f64) (x: [n] f64) : ([m] f64, [m] [n] f64) =
  (f x, tabulate n (jvp f x <-< idd n) |> transpose)
```

Functions wrapped with the `wrapj` function can now be used directly with the `dps.poly` function. Let's try it out in practice:

```
entry test_poly_jvp (sa_max:i64) : ([2] f64, bool, i64, i64, i64, f64) =
  let v0 = [0.5, 0.5]
  let ap = dps.default with sa_max = sa_max
  let (res, _, b, i, j, k, tol) =
    dps.poly ((wrapj <-< wrap2) bellman) v0 ap 0
  in (res, b, i, j, k, tol)

> test_poly_jvp 5i64

([0.7861513777574233f64, 0.6180339887498949f64], true, 5i64, 4i64, 1i64,
 1.1102230246251565e-16f64)
```

We see that we get the same results with `test_poly_jvp` as we get with `test_poly`.

Going parallel

The iterative approaches that the `dpsolver` functionality implements for finding fixpoints are inherently sequential, except from the matrix operations applied in the Newton-Kantorovich iterations (assuming a high-number of dimensions). Instead of parallelising the actual fixpoint resolution, we shall see how we can find many fixpoints in parallel, which is sometimes a useful approach for speeding up an application.

Following up on the task of finding intersection points between a circle and a simple quadratic equation, let us investigate how the x-dimension of the intersection points changes when the

circle radius increases.

We first parameterise the bellman equation over the radius of the circle:

```
let bellmanr (r:f64) (a:[2]f64) : ([2]f64, [2][2]f64) =  
  let f (a:[2]f64) = [f64.sqrt a[1], f64.sqrt(r**2-a[0]**2)]  
  let res = f a  
  let j = [[0, 1/(2*f64.sqrt a[1]) ],  
            [-a[0]/f64.sqrt(r**2-a[0]**2), 0 ]]  
  in (res, j)
```

We then create an entry point that implements an outer map over a call to `dps.poly` with varying radius:

```
let linspace (n: i64) (start: f64) (end: f64) : [n]f64 =  
  tabulate n (\i -> start + f64.i64 i * ((end-start)/f64.i64 n))  
  
entry test_polyr (n:i64) (sa_max:i64)  
  : (bool, i64, [n]f64, [n]f64) =  
  let ap = dps.default with sa_max = sa_max  
  let rs = linspace n 1 20  
  let res = map (\r -> let v0 = [0.5,0.5]  
                        let (res, _, b, i, j, _k, _tol) =  
                          dps.poly (bellmanr r) v0 ap 0  
                        in (r,res[0],b,i+j)) rs  
  let converged = reduce (&&) true (map (.2) res)  
  let xs = map (.1) res  
  let iterations = reduce (+) 0 (map (.3) res)  
  in (converged, iterations, rs, xs)
```

Here is a call to `test_polyr` with 4 different radius values (between 1 and 20) and an `sa_max` value of 3:

```
> test_polyr 4i64 3i64  
(true, 25i64, [1.0f64, 5.75f64, 10.5f64, 15.25f64],  
 [0.7861513777574233f64, 2.2960178985163853f64, 3.164158343195599f64, 3.841639561393954f64])
```

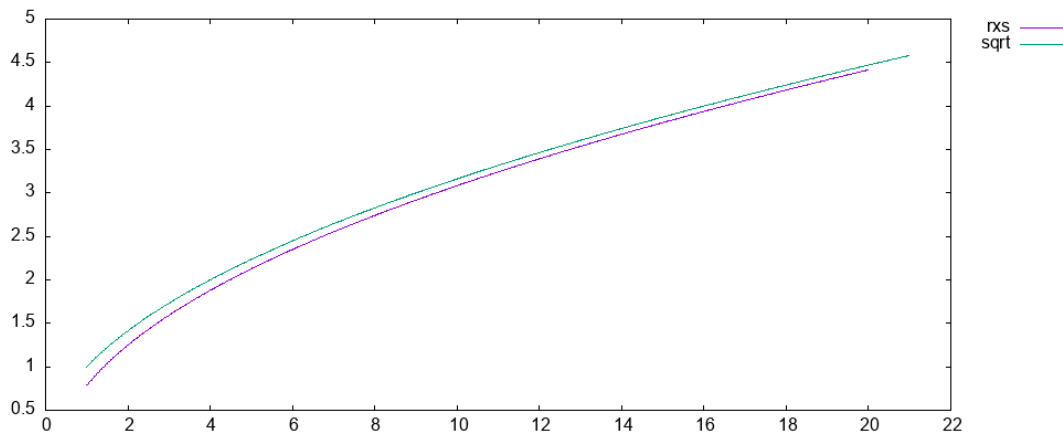
We can use the plot functionality of `iterate Futhark` to plot 1000 points relating radius values with associated found x -values (and compare it with a plot of the `sqrt`-function):

```
entry test_polyr_rxs (n:i64) (sa_max:i64)  
  : ([n]f64, [n]f64) =  
  test_polyr n sa_max |> (\(_,_,rs,xs) -> (rs,xs))  
  
let xys f n start end =  
  unzip (map (\x -> (x, f x)) (linspace n start end))  
  
entry sqrt_coords = xys f64.sqrt  
  
> :plot2d {rxs=test_polyr_rxs 1000i64 3i64,
```

```

sqrt=sqrt_coords 1000i64 1.0f64 21.0f64};
size: (400,1000)

```



A few single-dimensional examples

We now consider a single-dimensional case, for which we want to find the x for which $f(x) = \cos x$.

```

entry test_poly1d (sa_max : i64)
: ([1]f64, [1][1]f64, bool, i64, i64, i64, f64) =
let ap = dps.default with sa_max = sa_max
in dps.poly (\x -> ([f64.cos x[0]],
                    [[- f64.sin x[0]]]))
[0.7] ap 0

```

```
> test_poly1d 0i64
```

```

([0.7390851332151607f64], [[-0.6736120230211678f64]], true, 0i64, 3i64, 1i64,
0.0f64)

```

For another example, we want to compute $\sqrt{2}$ by finding the fixpoint to the equation $f(x) = \frac{1}{2}(x + \frac{2}{x})$.

```

entry test_sqrt (sa_max : i64)
: ([1]f64, [1][1]f64, bool, i64, i64, i64, f64) =
let ap = dps.default with sa_max = sa_max
in dps.poly (\x -> ( [ 0.5 * (x[0]+2/x[0]) ],
                    [[ 2*x[0] ]] )
) [1.4] ap 0

```

```
> test_sqrt 0i64
```

```

([1.414213562373095f64], [[2.828427124746191f64]], true, 0i64, 4i64, 1i64,
1.1102230246251565e-15f64)

```

Remarkably, in 4 steps we reach a fixpoint of 1.41421356237... with a tolerance of 1.11e-15.

Conclusion

We have seen how we can use the `dpsolver` library to solve multi-dimensional fixpoint equations. We have also seen how we can solve multiple problems in parallel using Futhark’s second-order array combinators.

References

- [1] Elsmann, M., Henriksen, T. and Oancea, C.E. 2018. *Parallel programming in Futhark*. Department of Computer Science, University of Copenhagen.
- [2] Henriksen, T., Serup, N.G.W., Elsmann, M., Henglein, F. and Oancea, C.E. 2017. Futhark: Purely functional GPU-programming with nested parallelism and in-place array updates. *Proceedings of the 38th ACM SIGPLAN conference on programming language design and implementation* (New York, NY, USA, 2017), 556–571.
- [3] Rust, J. 1996. Numerical dynamic programming in economics. Elsevier. 619–729.